Bachelor Thesis

Roman Sannikov

May 8, 2026

Contents

1 Introduction **2**

1.1 Literature review 3

2 Methodology **6**

2.1 Methods and technologies used 6

2.1.1 Project requirements 6

2.1.2 Technologies selection 6

2.1.3 Models 6

3 Results **7**

4 Conclusion **8**

5 Future work **9**

Reference list **10**

Chapter 1

Introduction

Over the past four decades, there was a major leap in robotics, especially in the field of quadrupeds. Many great papers developing the domain were published leading to significant advancements. The jump was from early days one-legged hopping machines to a four-legged "trotting" robots able to traverse through difficult terrain and obstacles, and even autonomous navigation in unstructured environments. Legged locomotion showed that it is one of the most versatile forms of mobility. But most of the research was done in labs, requiring substantial resources for prototyping and pushing the boundaries of the domain, which creates a gap between lab research and consumer-level accessibility.

The objective of this thesis is to address that gap by designing, constructing, and controlling a quadruped robot as an integrated project. The work covers the full development pipeline: mechanical design using modern CAD tools and 3D printing, selection and characterization of actuators, implementation of statically stable gaits based on established stability criteria, development of a control system, integration of computer vision, and validation through physical experiments. By following this end-to-end process, the thesis demonstrates how classical locomotion and computer vision principles can be realized on a budget-friendly platform.

A deliberate choice was made to use low-cost, commercially available components and rapid-prototyping techniques wherever possible. This decision was not driven by cost reduction alone; rather, it ensures that the resulting platform is reproducible. In this paper, I would like to review the research and practical achievements that have been done throughout the years

and combine it in a project, the cost of which will be less. The idea is to establish a base for low-cost quadruped robots to make them more affordable for day-to-day use in different fields of human lives.

1.1 Literature review

There is a number of "classical" papers that can be used to lay the foundation for this project. They vary from gaits and control to computer vision and successful real-world examples.

The most important thing for a quadruped robot is locomotion. The mathematical analysis in the 1968 "On the stability properties of quadruped creeping gaits" Mathematical Biosciences journal article shows, that there is "a unique optimum gait that maximizes static stability". The work explains that for static stability the center of gravity must be within the support polygon. This shape is formed by connecting all feet that touch the ground (contact points). Though, stability can be achived by dynamic means as well. The concept of dynamic stability is discussed in the 1984 "Experiments in Balance with a 3D One-Legged Hopping Machine", which explains the control algorithms: forward velocity, body altitude, and hopping height. In the recent years, compute became more accessible and the development of neural networks allowed to creation of simulated environments, in which ditigal versions of robots were put in and given a task to arrive from point A to B without being programmed to do so using Deep Reinforcement Learning (DRL). This approach was used in 2019 "Learning Agile and Dynamic Motor Skills for Legged Robots" for ANYmal robot. The researchers applied the real-world behavior of the robot's motors to the simulation that allowed the simulation to be more realistic and helped the robot learn imperfect nuances of the world.

While the theoretical foundations establish how quadruped locomotion works, several landmark systems demonstrate what is achievable with mature engineering. Boston Dynamics' BigDog (2008) was the first to demonstrate load-carrying dynamic locomotion on rough terrain, proving quadrupeds could operate in real-world conditions outside the lab. It was a direct descendant of the 1984 hopping machine experiments and it made evident that dynamic stability is the key to traversing challenging terrain. Similarly, MIT's Cheetah 3 (2018) showed that model-based control could achieve robust blind locomotion over obstacles. These systems represented the state-of-the-art capability of quadruped robots at their time. The projects were

sponsored by DARPA and MIT respectively and did not focus on cost reduction, but rather on pushing the boundaries of what quadruped robots could do.

Another key component of the project is computer vision. Generally, robots can be controlled by a human operator, which remains a very practical approach. However, computer vision is a significant step towards autonomy and high-level task execution. Vision is critical for navigation, obstacle avoidance, and interaction with the environment. For this project, I decided to use YOLO, a state-of-the-art object detection model, based on the "You Only Look Once" (2016) architecture, which was a game-changer in the field of real-time object detection. The model divides the image into a grid, and each grid cell is responsible for predicting bounding boxes, confidence score, and class probabilities of an object whose center falls into that cell. As described in the paper: "[The] network architecture is inspired by the GoogLeNet model for image classification... [It] has 24 convolutional layers followed by 2 fully connected layers. Instead of the inception modules used by GoogLeNet,... [it simply uses] 1×1 reduction layers followed by 3×3 convolutional layers...". The original YOLO uses the Sum-Squared Error (SSE) loss function between the predictions and the ground truth. While the architecture has evolved significantly in subsequent versions to utilize CSPDarknet, a high-performance convolutional neural network backbone, and advanced loss functions, the fundamental principle of grid-based single-pass detection remains the core of the model used in this project. Additionally, it is worth mentioning Object Detection with Discriminatively Trained Part Based Models (2010), which is a classic paper on training models to recognize objects based on their parts, because quadruped robots usually have limited field of view due to their height. Applying the logic of part-based modeling to train a model to recognize an object by its lower part is crucial for the robot to "understand" the environment.

While the architecture and approach to train the model are important, the datasets used for training enable the model to reach its full potential. In this project, I utilize two of the most significant benchmarks in the history of object detection: Common Objects in Context (COCO) (2014) and Stanford Dogs (2011). The COCO dataset established the standard for detecting objects in "natural" settings rather than perfectly centered studio shots. It contains over 330K images, with more than 2.5 million labeled instances across 91 object categories, including humans, dogs, cats, etc., easily recognizable by a 4-year-old. The Stanford Dogs dataset "...includes over 22,000 annotated images of dogs belonging to 120 species. Each image is

annotated with a bounding box and object class label.”

Chapter 2

Methodology

2.1 Methods and technologies used

2.1.1 Project requirements

2.1.2 Technologies selection

2.1.3 Models

Chapter 3

Results

Chapter 4

Conclusion

Chapter 5

Future work

Reference list

1. BigDog, the Rough-Terrain Quadruped Robot 2. MIT Cheetah 3: Design and Control of a Robust, Dynamic Quadruped Robot 3. On the stability properties of quadruped creeping gaits 4. Experiments in Balance with a 3D One-Legged Hopping Machine 5. Learning Agile and Dynamic Motor Skills for Legged Robots 6. You Only Look Once: Unified, Real-Time Object Detection 7. Object Detection with Discriminatively Trained Part-Based Models 8. Microsoft COCO: Common Objects in Context 9. Novel dataset for Fine-Grained Image Categorization: Stanford Dogs