Bachelor Thesis

Roman Sannikov

May 22, 2026

Contents

1 Introduction **2**

1.1 Literature review 3

2 Methodology **6**

2.1 Methods and technologies used 6

2.1.1 Project requirements 6

2.1.2 Technologies selection 7

2.1.3 Implementation 9

2.1.4 Models 12

3 Results **13**

4 Conclusion **14**

5 Future work **15**

Reference list **16**

Chapter 1

Introduction

Over the past four decades, there was a major leap in robotics, especially in the field of quadrupeds. Many great papers developing the domain were published leading to significant advancements. The jump was from early days one-legged hopping machines to a four-legged "trotting" robots able to traverse through difficult terrain and obstacles, and even autonomous navigation in unstructured environments. Legged locomotion showed that it is one of the most versatile forms of mobility. But most of the research was done in labs, requiring substantial resources for prototyping and pushing the boundaries of the domain, which creates a gap between lab research and consumer-level accessibility.

The objective of this thesis is to address that gap by designing, constructing, and controlling a quadruped robot as an integrated project. The work covers the full development pipeline: mechanical design using modern CAD tools and 3D printing, selection and characterization of actuators, implementation of statically stable gaits based on established stability criteria, development of a control system, integration of computer vision, and validation through physical experiments. By following this end-to-end process, the thesis demonstrates how classical locomotion and computer vision principles can be realized on a budget-friendly platform.

A deliberate choice was made to use low-cost, commercially available components and rapid-prototyping techniques wherever possible. This decision was not driven by cost reduction alone; rather, it ensures that the resulting platform is reproducible. In this paper, I would like to review the research and practical achievements that have been done throughout the years

and combine it in a project, the cost of which will be less. The idea is to establish a base for low-cost quadruped robots to make them more affordable for day-to-day use in different fields of human lives.

1.1 Literature review

There is a number of "classical" papers that can be used to lay the foundation for this project. They vary from gaits and control to computer vision and successful real-world examples.

The most important thing for a quadruped robot is locomotion. The mathematical analysis in the 1968 "On the stability properties of quadruped creeping gaits" Mathematical Biosciences journal article shows, that there is "a unique optimum gait that maximizes static stability". The work explains that for static stability the center of gravity must be within the support polygon. This shape is formed by connecting all feet that touch the ground (contact points). Though, stability can be achived by dynamic means as well. The concept of dynamic stability is discussed in the 1984 "Experiments in Balance with a 3D One-Legged Hopping Machine", which explains the control algorithms: forward velocity, body altitude, and hopping height. In the recent years, compute became more accessible and the development of neural networks allowed to creation of simulated environments, in which ditigal versions of robots were put in and given a task to arrive from point A to B without being programmed to do so using Deep Reinforcement Learning (DRL). This approach was used in 2019 "Learning Agile and Dynamic Motor Skills for Legged Robots" for ANYmal robot. The researchers applied the real-world behavior of the robot's motors to the simulation that allowed the simulation to be more realistic and helped the robot learn imperfect nuances of the world.

While the theoretical foundations establish how quadruped locomotion works, several landmark systems demonstrate what is achievable with mature engineering. Boston Dynamics' BigDog (2008) was the first to demonstrate load-carrying dynamic locomotion on rough terrain, proving quadrupeds could operate in real-world conditions outside the lab. It was a direct descendant of the 1984 hopping machine experiments and it made evident that dynamic stability is the key to traversing challenging terrain. Similarly, MIT's Cheetah 3 (2018) showed that model-based control could achieve robust blind locomotion over obstacles. These systems represented the state-of-the-art capability of quadruped robots at their time. The projects were

sponsored by DARPA and MIT respectively and did not focus on cost reduction, but rather on pushing the boundaries of what quadruped robots could do.

Another key component of the project is computer vision. Generally, robots can be controlled by a human operator, which remains a very practical approach. However, computer vision is a significant step towards autonomy and high-level task execution. Vision is critical for navigation, obstacle avoidance, and interaction with the environment. For this project, I decided to use YOLO, a state-of-the-art object detection model, based on the "You Only Look Once" (2016) architecture, which was a game-changer in the field of real-time object detection. The model divides the image into a grid, and each grid cell is responsible for predicting bounding boxes, confidence score, and class probabilities of an object whose center falls into that cell. As described in the paper: "[The] network architecture is inspired by the GoogLeNet model for image classification... [It] has 24 convolutional layers followed by 2 fully connected layers. Instead of the inception modules used by GoogLeNet,... [it simply uses] 1×1 reduction layers followed by 3×3 convolutional layers...". The original YOLO uses the Sum-Squared Error (SSE) loss function between the predictions and the ground truth. While the architecture has evolved significantly in subsequent versions to utilize CSPDarknet, a high-performance convolutional neural network backbone, and advanced loss functions, the fundamental principle of grid-based single-pass detection remains the core of the model used in this project. Additionally, it is worth mentioning Object Detection with Discriminatively Trained Part Based Models (2010), which is a classic paper on training models to recognize objects based on their parts, because quadruped robots usually have limited field of view due to their height. Applying the logic of part-based modeling to train a model to recognize an object by its lower part is crucial for the robot to "understand" the environment.

While the architecture and approach to train the model are important, the datasets used for training enable the model to reach its full potential. In this project, I utilize two of the most significant benchmarks in the history of object detection: Common Objects in Context (COCO) (2014) and Stanford Dogs (2011). The COCO dataset established the standard for detecting objects in "natural" settings rather than perfectly centered studio shots. It contains over 330K images, with more than 2.5 million labeled instances across 91 object categories, including humans, dogs, cats, etc., easily recognizable by a 4-year-old. The Stanford Dogs dataset "...includes over 22,000 annotated images of dogs belonging to 120 species. Each image is

annotated with a bounding box and object class label."

Chapter 2

Methodology

2.1 Methods and technologies used

Developing a functional quadruped robot on a limited budget requires navigating a wide landscape of hardware components, software frameworks, and design methodologies. While the robotics field offers a rich variety of actuators, computing platforms, sensors, and control strategies, the project constraints significantly narrow the range of viable options. Cost, reproducibility, and the mechanical demands of legged locomotion impose trade-offs at every stage of the design process, forcing a careful balance between performance, affordability, and ease of integration. The following subsections outline the requirements that shaped this design space and present the technologies and models selected to address them.

2.1.1 Project requirements

There are quite a lot of requirements and limitations in this project as it seeks to demonstrate a way to develop a low-cost quadruped robot. There are three big groups of limitations that define requirements. The first, and most significant, is cost. It influences the other group which is the hardware. The hardware, limited by cost, in turn constrains the software and algorithms. All together they create a really tight framework for the project, which is a good thing since it forces to creativity and finding new solutions.

My requirements for the project are the following:

- The robot should be able to walk on flat surfaces in all directions, i.e. forward, backwards, sideways, and turning in place. Flat surfaces were chosen here because rough terrain is too complex for the hardware the project uses.

- The robot should have 12 degrees of freedom (DOF) — where each joint that can move independently contributes one DOF — to allow for a variety of gaits and maneuvers. Each leg should have 3 DOF: hip abduction/adduction (side-to-side movement), hip flexion/extension (forward-backward movement), and knee flexion/extension.

- It should have computer vision capabilities to detect and recognize objects. For the sake of demonstration, I chose dogs and humans as the target objects because they represent the primary agents (humans) and subjects (dogs) the robot is expected to encounter and interact with.

- The robot should be easy to assemble and reproduce by others. The components should be available, affordable, and designed with swappability in mind to allow for future upgrades and/or repairs.

On the limitation side, the cost of the whole project should not exceed 300 EUR, which is quite limiting for the result the project wants to achieve.

The hardware and software choices are affected by this limitation and the technology selection will be discussed in the next section.

2.1.2 Technologies selection

Everything starts with the model. Nowadays, 3D printers are quite accessible for people making them a great choice for the project since it allows producing custom elements perfectly fitting and satisfying the requirements and at the same time being easily reproducible by the others. At my disposal were Bambu Lab printers A1 mini, X1C Carbon and H2D series. I tried them all and settled on the last one since it has two nozzles, allowing the use of two types of filament at the same time: PLA and PETG. PLA is a great general material for the robot itself, while PETG was used for supporting elements. It is possible to use only PLA, and I have tried that, but PETG makes life easier since it does not stick to PLA, making it a perfect choice for supports. The design of the robot was done in Onshape, a free cloud-based CAD web software with enough

functionality for the project. For the slicer — software that converts 3D models into G-code instructions for the printer, specifying toolpaths, layer heights, and print parameters — Bambu Lab's Bambu Studio was used as it perfectly works with the printers and has the sufficient functionality for the project. I have not created the initial model by myself, as it would take much more time, but rather found a suitable one with Creative Commons (4.0 International License) on the Internet made by Tom Knox and modified it to fit the requirements. The requirements not met by the initial model were 12 DOF and ease of assembly. What the original model did provide was a solid base and design I wanted, 3 DOF per front leg, and a place for the camera module.

What made the original 3D model hard to assemble was the fact that it consisted of 149 unique parts connected by different types of screws. The most problematic part was the leg. All four legs are the same just mirrored, which made redesigning much easier. To reduce complexity of assembly, I redesigned legs and the body to be print-in-place, meaning ready to use parts without the need of additional assembly. I also reduced the number of unique screws in the project to just 1 type, which is M3 of two sizes.

For the actuators, I chose WS-MG996R digital servos with metal gears, which are widely used in hobby robotics, 3 per leg, which gives 12 DOF in total. These servos are really affordable and deliver 11 $kg \cdot cm$ stall torque with a speed of 0.17–0.14 seconds per 60°, operating at 4.8V to 7.2V DC. On the other hand, cheaper actuators mean low precision, feedback, and torque, which are necessary for this type of project. Considering all pros and cons, these are the bare minimum for the project.

For the main controller board, I chose Raspberry Pi 5 4GB. It is a nice little computer powerful enough for controlling servos and running small object recognition models. A big advantage of this computer is that it runs Raspberry Pi OS, which is a Linux based operating system, allowing fast prototyping and running existing programs and tools. The Raspberry Pi 5 offers significant improvements over its predecessor, including faster GPIO and better overall performance, making it well-suited for this application. The Raspberry Pi does not have enough GPIO pins to control all 12 servos, so I went with PCA9685, a 16-channel, 12-bit PWM (Pulse Width Modulation) driver chip. It uses I2C communication protocol and takes just 4 Raspberry Pi GPIO pins: VCC 3V, SDA, SCL, and GND. The servos consumption can spike up to 2.5A per servo, so in total it is 30A, which is way more than the Raspberry Pi can provide. To solve

this problem, I use an external power supply, which is 4 18650 2500 mAh Lithium-ion batteries able to produce enough current for the servos. Connected in a 2S2P configuration (two series pairs in parallel), they provide 7.4V with a total capacity of 5000 mAh. With that the potential difference is larger than what the servos and Raspberry Pi can handle, so I use buck converters to step down the voltage to 7V for the servos and 5V for the Raspberry Pi.

For the computer vision part, I use a 5-megapixel Raspberry Pi Camera Module v1.3, which features an OmniVision OV5647 Color CMOS sensor. The camera has a 3.6 mm focal length lens with a $54^\circ \times 41^\circ$ angle of view, providing a horizontal field of view roughly equal to the distance from the subject at typical operating ranges. The camera uses a CSI ribbon cable to connect to the high-speed, dual-purpose MIPI port on the Raspberry Pi, offering lower latency and CPU overhead compared to USB cameras.

On the software side, I use Python as the main programming language. It is a great choice for rapid prototyping and has a huge ecosystem of libraries for robotics, computer vision, and machine learning. For controlling the servos, I use the Adafruit PCA9685 library, which provides a simple interface for sending PWM (Pulse Width Modulation) signals to the servos. PWM varies the duty cycle of a square wave to control the average voltage delivered, with pulse widths typically ranging from 1 to 2 milliseconds corresponding to 0 to 180 degrees of servo rotation. For object detection, I use the Ultralytics YOLO library, which offers a streamlined interface for YOLO-based inference and integrates well with the Raspberry Pi platform.

2.1.3 Implementation

The implementation started by downloading the original model and printing the most important part — the leg. The leg model consisted of 7 parts, which took the printer around 3 hours. After, I assembled the leg, attached the hip to a servo and connected the servo directly to the Raspberry Pi with jumper wires. Printing it and testing that the leg works was a crucial step that helped me to verify that the model is capable of moving and having the physical model helped me understand what needed to be changed. Later, I started redesigning the leg to be print-in-place. First of all, it is important to understand the principles of such design. The most important idea is that there should be no overhanging parts, meaning that all the parts should be supported by the print itself. To achieve that, the maximum overhang angle should not exceed 45 degrees, beyond which support material would be required. Another thing to

keep in mind is the clearance between the moving parts. Through experimentation, I found that 0.65 mm provides the best balance between smooth motion and minimal play, so this value was used for most clearances in the model. The height of every print was standard — 0.2 mm and the support parameter was set for hanging parts like curves of the legs and body "ribs". The final leg consists of 6 elements: motor extension, extension link, shoulder disk, femur, secondary link, and tibia, printed as one module. After removing it from the printer bed and detaching the supports, it is movable and ready to be installed. To get to the final design, I went through around twenty iterations of the leg and a couple of iterations of the body. Several design challenges emerged during these iterations and to achieve the requirements while using the same design, I had to break the basic principles. For example, a support blocker is used for the leg to block the support from being generated in the motor extension part, creating a bridging surface that spans a gap without support. Another example, showing the challenge of designing print-in-place models, is the knee. The knee joint employs a cylindrical pivot pin that is integrated into the tibia (lower leg), with the femur (upper leg) captured between the two halves of the tibia joint — forming a sandwich-like structure. This arrangement creates overhangs on both the top and bottom surfaces of the femur where it interfaces with the tibia, requiring support material to be generated in two places: between the upper tibia surface and the femur, and between the femur and the lower tibia surface. The overhang surface is large and cannot be recovered if drooped and melted to the other part, thus support was necessary, so unfortunately, each knee has a bit of PETG support in it. PLA filament can also be used as an alternative to PETG for support, requiring breaking it by moving the knee joint after printing. Beyond the main structural parts, the model contains eight servo holders that are printed separately. Each holder weighs 5 grams and does not require support material. The holders attach to the servos with screws, while the connection to the body uses integrated pins that slot into corresponding holes, eliminating the need for additional fasteners. In total, printing the complete model requires approximately 20.5 hours and 490 grams of filament. The body takes 7 hours and 154 grams, each of the four legs requires 3 hours and 74 grams, and the eight servo holders together take 1.5 hours and 40 grams.

When the whole model was stable and ready, I started assembling it all together. With the experience I gained, it would take around one hour to assemble everything. Connecting PCA9685 to the Raspberry Pi was a straightforward process. The control pin connection layout is the following: VCC to 3.3V, GND to GND, SDA to GPIO 2 (SDA), and SCL to GPIO 3 (SCL).

PCA9685 does not power the servos from the controller, instead there is a separate V+ and GND terminals. Those are connected to an external power supply. While testing I used a lab power supply, set to 6V and 6A, which was enough for testing purposes. In order to enable I2C communication interface on the Raspberry Pi, it needed to be enabled through the Raspberry Pi Configuration tool, which can be accessed through `sudo raspi-config` under Interface Options, I2C. After that, I2C utilities should be installed by running `sudo apt install i2c-tools`. The connectivity can be verified by running `i2cdetect -y 1`, which should show the PCA9685 at address $0x40$ (hexadecimal). The address is hardcoded to $0x40$ by the manufacturer, as the default base address for this chip. In the I2C protocol, each device requires a unique 7-bit address for the master controller to identify it. The PCA9685 uses a partially configurable address scheme where the upper bits are fixed internally (binary 1000000, corresponding to hexadecimal $0x40$), while the lower six bits can be modified via external address pins labeled A^0 through A^5. This design allows up to 62 PCA9685 devices to coexist on the same bus by changing the pin configurations, enabling expansion of PWM outputs without requiring additional communication channels. For instance, setting pin A^0 high while keeping others low changes the address from $0x40$ to $0x41$. This flexible addressing makes the PCA9685 well-suited for robotics applications requiring numerous servo motors. For this project, I am not changing the default address, so it remains at $0x40$. The other pins are 16 channels for PWM output, which are connected to the servos. Each channel consists of three pins: V+ (power), GND (ground), and OUT (signal). I grouped them by leg in this order: hip extension, knee flexion, and hip abduction, so it is easier to understand which channel controls which joint. The PCA9685 board exposes sockets for the servos, so it is convenient to connect them without soldering, however, the board itself exposes the I2C pins, while Raspberry Pi exposes its GPIO pins, so I had to solder wires to connect them together. Each servo has 180° range of motion, so before attaching servos to the robot, I set: hip to 90°, knee to 60°, and abduction to 95°, to their home position. These values are hardcoded in the code and all movements are relative to them. Then, the servos are attached as following: the hip abduction servo is screwed vertically to its slot in the body and a dual-arm horn is attached to the gear, after that the hip extension and knee flexion servos are put together and secured with servo holders on the sides, then another dual-arm horn is attached to the knee flexion servo, while a round disc horn is screwed to the back side of the femur. Ultimately, the horns are connected to the hip and motor extension respectively. The same procedure is repeated for all four legs. After that, the

electronics are installed inside the body and the Raspberry Pi camera module is mounted to the front camera slot. With the mechanical assembly complete, the control software can command movement.

A step consists of two phases: the swing phase, where the leg is parabolically lifted and moved forward, and the stance phase, where the foot drives backward relative to the chassis, generating the ground reaction force that pushes the robot's body forward. A gait is the periodic pattern of these foot movements. To translate desired foot positions into servo angles, the system uses a lookup table approach rather than inverse kinematics. This method was chosen because the mechanical coupling in the leg design — where hip movement affects knee angle — makes standard analytical solutions impractical. The lookup table contains pre-calibrated mappings between joint angles and foot positions, created by setting different servo angles and measuring the resulting foot positions. When a target position is requested, the system identifies the nearest calibration points and interpolates between them using inverse distance weighting, providing smooth transitions even with sparse calibration data. Each servo is individually calibrated to account for mechanical variations and assembly differences between left and right legs. A direction multiplier ensures that positive commands produce consistent physical movements across all legs, while center offsets compensate for variations in the home position. Three distinct gaits are implemented: a crawl gait for statically stable locomotion, where legs move sequentially in a diagonal pattern; a trot gait for faster movement, where diagonal leg pairs move simultaneously; and a turning gait, where left and right legs move in opposite directions to generate rotational torque. Servo movements are interpolated over time with configurable speed, ensuring smooth motion and preventing sudden movements that could upset the robot's balance. All movements are defined relative to a home position, which serves as the reference point for calibration and initialization. Since the lookup table was measured for the left legs, the right legs use mirrored values to maintain symmetry in movement.

2.1.4 Models

Chapter 3

Results

Chapter 4

Conclusion

Chapter 5

Future work

Bibliography

[1] Raibert, M., Blankespoor, K., Nelson, G., & Playter, R. (2008). BigDog, the rough-terrain quadruped robot. *Proceedings of the 17th IFAC World Congress*, 10822–10825.

[2] Katz, B., Di Carlo, J., & Kim, S. (2019). MIT Cheetah 3: Design and control of a robust, dynamic quadruped robot. *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, 2245–2251.

[3] McGhee, R. B., & Frank, A. A. (1968). On the stability properties of quadruped creeping gaits. *Mathematical Biosciences*, 3(3–4), 331–351.

[4] Raibert, M. H. (1984). Experiments in balance with a 3D one-legged hopping machine. *The International Journal of Robotics Research*, 3(2), 75–92.

[5] Hwangbo, J., Lee, J., Dosovitskiy, A., Bellicoso, D., Tsounis, V., Koltun, V., & Hutter, M. (2019). Learning agile and dynamic motor skills for legged robots. *Science Robotics*, 4(26), eaau5872.

[6] Redmon, J., Divvala, S., Girshick, R., & Farhadi, A. (2016). You only look once: Unified, real-time object detection. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 779–788.

[7] Felzenszwalb, P. F., Girshick, R. B., McAllester, D., & Ramanan, D. (2010). Object detection with discriminatively trained part-based models. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 32(9), 1627–1645.

[8] Lin, T. Y., Maire, M., Belongie, S., Hays, J., Perona, P., Ramanan, D., Dollár, P., & Zitnick, C. L. (2014). Microsoft COCO: Common objects in context. *Proceedings of the*

European Conference on Computer Vision (ECCV), 740–755.

[9] Khosla, A., Jayadevaprakash, N., Yao, B., & Fei-Fei, L. (2011). Novel dataset for fine-grained image categorization. *First Workshop on Fine-Grained Visual Categorization, IEEE Conference on Computer Vision and Pattern Recognition*.

[10] Knox, T. (n.d.). Quadruped I. Retrieved from `https://www.printables.com/model/1443108-quadruped-i`